

Vrushali_new

December 17, 2024

```
[1]: spark = SparkSession.builder.appName("NYPD").getOrCreate()
      spark.sparkContext.setLogLevel("ERROR")
```

24/12/17 01:11:53 WARN SparkSession: Using an existing Spark session; only runtime SQL configurations will take effect.

```
[2]: import pyspark.sql.functions as F
      import matplotlib.pyplot as plt
      import pandas as pd
      import folium
      from folium.plugins import HeatMap
```

```
[3]: #https://storage.cloud.google.com/nypdbucket/notebooks/jupyter/nypd_dataset.csv
      csv_file_path = "gs://nypdbucket/notebooks/jupyter/nypd_dataset.csv"
```

```
# Read the CSV file into a DataFrame
df = (spark.read.format("csv")
      .option("header", "true") # If the CSV has a header row
      .option("inferSchema", "true") # Automatically infer the schema
      .load(csv_file_path))
```

```
# Show the first few rows of the DataFrame
df.show(5)
```

[Stage 2:> (0 + 1) / 1]

```
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
----+-----+-----+
|CAD_EVNT_ID|CREATE_DATE|INCIDENT_DATE|INCIDENT_TIME|NYPD_PCT_CD|BORO_NM|
|PATRL_BORO_NM|GEO_CD_X|GEO_CD_Y|RADIO_CODE|TYP_DESC|CIP_JOBS|
|ADD_TS|DISP_TS|ARRIVD_TS|CLOSNG_TS|Latitude|
|Longitude|
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
```

```

-----+-----+-----+
| 99842231| 01/01/2024| 12/31/2023|2024-12-17 22:11:38| 45| BRONX|
PATROL BORO BRONX| 1031438| 249344| 52D6| DISPUTE: FAMILY| Non
CIP|01/01/2024 12:02:...|01/01/2024 12:02:...| NULL|01/01/2024
12:13:...|40.850949|-73.829434|
| 99842388| 01/01/2024| 12/31/2023|2024-12-17 22:19:46| 110|
QUEENS|PATROL BORO QUEEN...| 1022087| 208229| 52D6| DISPUTE: FAMILY|
Non CIP|01/01/2024 12:09:...|01/01/2024 12:11:...|
NULL|01/01/2024 12:57:...|40.738144|-73.863466|
| 99842587| 01/01/2024| 12/31/2023|2024-12-17 22:31:06| 108|
QUEENS|PATROL BORO QUEEN...| 1007298| 209993| 53S|VEHICLE ACCIDENT:...|
Non CIP|01/01/2024 12:01:...|01/01/2024 12:02:...|01/01/2024
01:19:...|01/01/2024 01:20:...|40.743037|-73.916826|
| 99843964| 01/01/2024| 12/31/2023|2024-12-17 23:53:22| 114|
QUEENS|PATROL BORO QUEEN...| 1002279| 222019| 34K1|ASSAULT (IN
PROGR...|Critical|01/01/2024 12:06:...|01/01/2024 12:07:...|01/01/2024
12:19:...|01/01/2024 01:03:...|40.776057|-73.934906|
| 99844026| 01/01/2024| 12/31/2023|2024-12-17 23:57:38|
66|BROOKLYN|PATROL BORO BKLYN...| 987908| 174328| 11C4|ALARMS:
COMMERCIA...| Non CIP|01/01/2024 12:04:...|01/01/2024 01:45:...|
NULL|01/01/2024 02:23:...|40.645174| -73.98682|
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+
only showing top 5 rows

```

```
[4]: # Print the schema of the DataFrame
df.printSchema()
```

```

root
|-- CAD_EVNT_ID: integer (nullable = true)
|-- CREATE_DATE: string (nullable = true)
|-- INCIDENT_DATE: string (nullable = true)
|-- INCIDENT_TIME: timestamp (nullable = true)
|-- NYPD_PCT_CD: integer (nullable = true)
|-- BORO_NM: string (nullable = true)
|-- PATRL_BORO_NM: string (nullable = true)
|-- GEO_CD_X: integer (nullable = true)
|-- GEO_CD_Y: integer (nullable = true)
|-- RADIO_CODE: string (nullable = true)
|-- TYP_DESC: string (nullable = true)
|-- CIP_JOBS: string (nullable = true)
|-- ADD_TS: string (nullable = true)
|-- DISP_TS: string (nullable = true)
|-- ARRIVD_TS: string (nullable = true)

```

```

|-- CLOSNG_TS: string (nullable = true)
|-- Latitude: double (nullable = true)
|-- Longitude: double (nullable = true)

```

```
[5]: df.describe().show()
```

24/12/17 01:12:15 WARN SparkStringUtils: Truncated the string representation of a plan since it was too large. This behavior can be adjusted by setting 'spark.sql.debug.maxToStringFields'.

[Stage 5:> (0 + 1) / 1]

```

+-----+-----+-----+-----+-----+-----+
|summary|      CAD_EVNT_ID|CREATE_DATE|INCIDENT_DATE|      NYPD_PCT_CD|
BORO_NM|      PATRL_BORO_NM|      GEO_CD_X|      GEO_CD_Y|
RADIO_CODE|      TYP_DESC|CIP_JOBS|      ADD_TS|
DISP_TS|      ARRIVD_TS|      CLOSNG_TS|      Latitude|
Longitude|
+-----+-----+-----+-----+-----+-----+
| count|      5430525|      5430525|      5430525|      5430524|
5430525|      5430525|      5430525|      5430525|
5430525|      5430525| 5430525|      5430525|      5430525|
4292498|      5430492|      5430525|      5430525|
| mean|1.031539814432809E8|      NULL|      NULL| 61.0436114452307|
NULL|      NULL|1003719.7164797143|207219.3408618872|
4.88391838826703E8|      NULL|      NULL|      NULL|
NULL|      NULL|      NULL| 40.73540911380572|
-73.92972289303498|
| stddev| 1915287.3223354751|      NULL|      NULL|34.59641421997529|
NULL|
NULL|20214.115777641575|29594.63603927419|5.062920631799014E9|
NULL|      NULL|      NULL|      NULL|      NULL|
NULL|0.08123306598960091|0.07290432332565411|
| min|      99842231| 01/01/2024| 01/01/2024|      0|
(null)|      (null)|      913411|      121022|
001|10-53 NO RMP REQU...|Critical|01/01/2024 01:00:...|01/01/2024
01:00:...|01/01/2024 01:00:...|01/01/2024 01:00:...|      40.498596|
-74.254743|
| max|      106487582| 09/30/2024| 12/31/2023|      123|STATEN
ISLAND|PATROL BORO STATE...|      1067305|      272307|
DEPTOW|      YOUTH HOME VISIT| Serious|09/30/2024 12:59:...|10/01/2024
12:53:...|10/01/2024 12:59:...|10/01/2024 12:59:...|      40.914065|
-73.700291|

```

```

+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+

```

1 Data Cleaning

```
[6]: print("Duplicate Rows:", df.count() - df.distinct().count())
```

```
[Stage 11:=====> (12 + 1) / 13]
```

```
Duplicate Rows: 0
```

```
[7]: arr_null_count = df.filter(F.col("ARRIVD_TS").isNull()).count()
clo_null_count = df.filter(F.col("CLOSNG_TS").isNull()).count()
pct_null_count = df.filter(F.col("NYPD_PCT_CD").isNull()).count()

print(f"Number of rows with NULL values in ARRIVD_TS: {arr_null_count}")
print(f"Number of rows with NULL values in CLOSNG_TS: {clo_null_count}")
print(f"Number of rows with NULL values in NYPD_PCT_CD: {pct_null_count}")
```

```
[Stage 21:=====> (8 + 4) / 12]
```

```
Number of rows with NULL values in ARRIVD_TS: 1138027
```

```
Number of rows with NULL values in CLOSNG_TS: 33
```

```
Number of rows with NULL values in NYPD_PCT_CD: 1
```

```
[8]: # Drop rows where ARRIVD_TS or CLOSNG_TS is NULL
df_cleaned = df.filter((F.col("ARRIVD_TS").isNotNull()) & (F.col("CLOSNG_TS").
    ↪isNotNull()))

# Show counts after dropping rows
print("Total rows after dropping NULL values in ARRIVD_TS and CLOSNG_TS:",
    ↪df_cleaned.count())

# Verify NULLs again
df_cleaned.select(
    F.col("ARRIVD_TS").isNull().alias("ARRIVD_TS_NULL"),
    F.col("CLOSNG_TS").isNull().alias("CLOSNG_TS_NULL")
).show(5)
```

Total rows after dropping NULL values in ARRIVD_TS and CLOSNG_TS: 4292472

```
+-----+-----+
|ARRIVD_TS_NULL|CLOSNG_TS_NULL|
+-----+-----+
|           false|           false|
|           false|           false|
|           false|           false|
|           false|           false|
|           false|           false|
+-----+-----+
```

only showing top 5 rows

```
[9]: # Fill NULL value in NYPD_PCT_CD with 'UNKNOWN'
df_cleaned = df_cleaned.fillna({"NYPD_PCT_CD": "UNKNOWN"})

# Verify NULLs in NYPD_PCT_CD
print("NULL values in NYPD_PCT_CD after filling:")
df_cleaned.filter(F.col("NYPD_PCT_CD").isNull()).show()
```

NULL values in NYPD_PCT_CD after filling:

[Stage 30:=====> (3 + 4) / 7]

```
+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+
|CAD_EVNT_ID|CREATE_DATE|INCIDENT_DATE|INCIDENT_TIME|NYPD_PCT_CD|BORO_NM|PATRL_B
ORO_NM|GEO_CD_X|GEO_CD_Y|RADIO_CODE|TYP_DESC|CIP_JOBS|ADD_TS|DISP_TS|ARRIVD_TS|C
LOSNG_TS|Latitude|Longitude|
+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+
```

[Stage 30:=====> (6 + 1) / 7]

```
[10]: # Verify NULL counts for all relevant columns
df_cleaned.select(
    F.count(F.when(F.col("ARRIVD_TS").isNull(), 1)).
    ↪alias("ARRIVD_TS_NULL_COUNT"),
    F.count(F.when(F.col("CLOSNG_TS").isNull(), 1)).
    ↪alias("CLOSNG_TS_NULL_COUNT"),
    F.count(F.when(F.col("NYPD_PCT_CD").isNull(), 1)).
    ↪alias("NYPD_PCT_CD_NULL_COUNT")
)
```

```
).show()
```

```
[Stage 31:=====> (11 + 1) / 12]
```

```
+-----+-----+-----+
|ARRIVD_TS_NULL_COUNT|CLOSNG_TS_NULL_COUNT|NYPD_PCT_CD_NULL_COUNT|
+-----+-----+-----+
|              0|              0|              0|
+-----+-----+-----+
```

```
[11]: # Drop unnecessary columns
df_cleaned = df_cleaned.drop("CAD_EVNT_ID", "PATRL_BORO_NM")

# Verify the columns after dropping
print("Columns after dropping CAD_EVNT_ID and PATRL_BORO_NM:")
df_cleaned.printSchema()
```

Columns after dropping CAD_EVNT_ID and PATRL_BORO_NM:

root

```
-- CREATE_DATE: string (nullable = true)
-- INCIDENT_DATE: string (nullable = true)
-- INCIDENT_TIME: timestamp (nullable = true)
-- NYPD_PCT_CD: integer (nullable = true)
-- BORO_NM: string (nullable = true)
-- GEO_CD_X: integer (nullable = true)
-- GEO_CD_Y: integer (nullable = true)
-- RADIO_CODE: string (nullable = true)
-- TYP_DESC: string (nullable = true)
-- CIP_JOBS: string (nullable = true)
-- ADD_TS: string (nullable = true)
-- DISP_TS: string (nullable = true)
-- ARRIVD_TS: string (nullable = true)
-- CLOSNG_TS: string (nullable = true)
-- Latitude: double (nullable = true)
-- Longitude: double (nullable = true)
```

```
[12]: from pyspark.sql.functions import to_timestamp

# Convert columns to timestamp format
df_cleaned = df_cleaned.withColumn("ADD_TS", to_timestamp("ADD_TS", "MM/dd/yyyy_
↳hh:mm:ss a")) \
                        .withColumn("DISP_TS", to_timestamp("DISP_TS", "MM/dd/
↳yyyy hh:mm:ss a")) \
                        .withColumn("ARRIVD_TS", to_timestamp("ARRIVD_TS", "MM/
↳dd/yyyy hh:mm:ss a")) \
```

```

        .withColumn("CLOSNG_TS", to_timestamp("CLOSNG_TS", "MM/
        ↪dd/yyyy hh:mm:ss a"))

# Verify the schema to ensure correct types
print("Schema after converting to timestamp:")
df_cleaned.printSchema()

# Show sample rows
df_cleaned.select("ADD_TS", "DISP_TS", "ARRIVD_TS", "CLOSNG_TS").show(5,
        ↪truncate=False)

```

Schema after converting to timestamp:

```

root
 |-- CREATE_DATE: string (nullable = true)
 |-- INCIDENT_DATE: string (nullable = true)
 |-- INCIDENT_TIME: timestamp (nullable = true)
 |-- NYPD_PCT_CD: integer (nullable = true)
 |-- BORO_NM: string (nullable = true)
 |-- GEO_CD_X: integer (nullable = true)
 |-- GEO_CD_Y: integer (nullable = true)
 |-- RADIO_CODE: string (nullable = true)
 |-- TYP_DESC: string (nullable = true)
 |-- CIP_JOBS: string (nullable = true)
 |-- ADD_TS: timestamp (nullable = true)
 |-- DISP_TS: timestamp (nullable = true)
 |-- ARRIVD_TS: timestamp (nullable = true)
 |-- CLOSNG_TS: timestamp (nullable = true)
 |-- Latitude: double (nullable = true)
 |-- Longitude: double (nullable = true)

```

```

+-----+-----+-----+-----+
+
|ADD_TS          |DISP_TS          |ARRIVD_TS        |CLOSNG_TS
|
+-----+-----+-----+-----+
+
|2024-01-01 00:01:21|2024-01-01 00:02:19|2024-01-01 01:19:58|2024-01-01
01:20:02|
|2024-01-01 00:06:11|2024-01-01 00:07:19|2024-01-01 00:19:27|2024-01-01
01:03:22|
|2024-01-01 00:04:51|2024-01-01 00:09:21|2024-01-01 00:15:11|2024-01-01
00:56:56|
|2024-01-01 00:04:57|2024-01-01 00:12:08|2024-01-01 00:29:16|2024-01-01
00:29:53|
|2024-01-01 00:00:07|2024-01-01 00:00:07|2024-01-01 00:00:07|2024-01-01
00:30:23|
+-----+-----+-----+-----+
+

```

only showing top 5 rows

```
[13]: from pyspark.sql.functions import unix_timestamp, round

# Calculate correct response time differences in minutes
df_cleaned = df_cleaned.withColumn("dispatch_time",
    round((unix_timestamp("DISP_TS") - unix_timestamp("ADD_TS")) / 60, 2)) \
    .withColumn("arrival_time",
    round((unix_timestamp("ARRIVD_TS") - unix_timestamp("DISP_TS")) / 60, 2)) \
    .withColumn("total_response_time",
    round((unix_timestamp("ARRIVD_TS") - unix_timestamp("ADD_TS")) / 60, 2))

# Show the corrected metrics
df_cleaned.select("ADD_TS", "DISP_TS", "ARRIVD_TS", "dispatch_time",
    "arrival_time", "total_response_time").show(10)
```

```
+-----+-----+-----+-----+-----+
-----+-----+
|          ADD_TS|          DISP_TS|
ARRIVD_TS|dispatch_time|arrival_time|total_response_time|
+-----+-----+-----+-----+-----+
-----+-----+
|2024-01-01 00:01:21|2024-01-01 00:02:19|2024-01-01 01:19:58|          0.97|
77.65|          78.62|
|2024-01-01 00:06:11|2024-01-01 00:07:19|2024-01-01 00:19:27|          1.13|
12.13|          13.27|
|2024-01-01 00:04:51|2024-01-01 00:09:21|2024-01-01 00:15:11|          4.5|
5.83|          10.33|
|2024-01-01 00:04:57|2024-01-01 00:12:08|2024-01-01 00:29:16|          7.18|
17.13|          24.32|
|2024-01-01 00:00:07|2024-01-01 00:00:07|2024-01-01 00:00:07|          0.0|
0.0|          0.0|
|2024-01-01 00:00:14|2024-01-01 00:08:24|2024-01-01 00:36:32|          8.17|
28.13|          36.3|
|2024-01-01 00:00:25|2024-01-01 00:00:25|2024-01-01 00:00:25|          0.0|
0.0|          0.0|
|2024-01-01 00:00:35|2024-01-01 00:00:35|2024-01-01 00:00:35|          0.0|
0.0|          0.0|
|2024-01-01 00:05:03|2024-01-01 00:15:06|2024-01-01 00:42:21|         10.05|
27.25|          37.3|
|2024-01-01 00:00:51|2024-01-01 00:22:17|2024-01-01 00:30:42|         21.43|
8.42|          29.85|
+-----+-----+-----+-----+-----+
-----+-----+
only showing top 10 rows
```



```
[14]: from pyspark.sql.functions import col

# Identify outliers for dispatch_time and arrival_time using 95th percentile
dispatch_quantile = df_cleaned.approxQuantile("dispatch_time", [0.95], 0.0)[0]
arrival_quantile = df_cleaned.approxQuantile("arrival_time", [0.95], 0.0)[0]

print(f"95th percentile of dispatch_time: {dispatch_quantile}")
print(f"95th percentile of arrival_time: {arrival_quantile}")

# Filter out extreme values for dispatch_time and arrival_time
df_cleaned = df_cleaned.filter(
    (col("dispatch_time") > 0) & (col("dispatch_time") <= dispatch_quantile) &
    (col("arrival_time") > 0) & (col("arrival_time") <= arrival_quantile)
)

# Verify the filtered dataset
df_cleaned.describe("dispatch_time", "arrival_time").show()
```

95th percentile of dispatch_time: 20.22

95th percentile of arrival_time: 101.27

[Stage 40:=====> (11 + 1) / 12]

summary	dispatch_time	arrival_time
count	1519219	1519219
mean	3.215696466408669	16.94440944327612
stddev	4.220197489145211	22.136328578888538
min	0.02	0.02
max	20.22	101.27

```
[15]: # Filter rows with negative response times or any invalid values
df_cleaned = df_cleaned.filter(
    (col("dispatch_time") > 0) &
    (col("arrival_time") > 0) &
    (col("total_response_time") > 0)
)

# Verify cleaned data
print("After filtering invalid values:")
df_cleaned.select("dispatch_time", "arrival_time", "total_response_time").
    describe().show()
```

After filtering invalid values:

[Stage 43:=====> (9 + 3) / 12]

summary	dispatch_time	arrival_time	total_response_time
count	1519219	1519219	1519219
mean	3.215696466408669	16.94440944327612	20.160261772661084
stddev	4.220197489145211	22.136328578888538	23.61289080671234
min	0.02	0.02	0.03
max	20.22	101.27	121.35

```
[16]: from pyspark.ml.feature import StandardScaler, VectorAssembler

# Assemble features for scaling
assembler = VectorAssembler(inputCols=["dispatch_time", "arrival_time"],
    ↪outputCol="raw_features")
df_scaled = assembler.transform(df_cleaned)

# Apply StandardScaler
scaler = StandardScaler(inputCol="raw_features", outputCol="features",
    ↪withStd=True, withMean=True)
scaler_model = scaler.fit(df_scaled)
df_final = scaler_model.transform(df_scaled).select("features",
    ↪"total_response_time")

# Show scaled data
print("Scaled Features:")
df_final.show(5, truncate=False)
```

Scaled Features:

features	total_response_time
[-0.5321306579097865, 2.7423513497455883]	78.62
[-0.4942177402296219, -0.21748906672196322]	13.27
[0.3043230884088464, -0.5020891067669201]	10.33
[0.9393644595516045, 0.008383980932764353]	24.32
[1.1739506376976234, 0.5053046857731651]	36.3

only showing top 5 rows

```
[17]: # Combine all partitions into one
df_cleaned = df_cleaned.coalesce(1)

# Specify the output folder path
output_folder = "gs://nypdbucket/notebooks/jupyter/cleaned_dataset"

# Write the cleaned dataset as a single CSV file
df_cleaned.write.mode("overwrite").option("header", "true").csv(output_folder)

print("Final cleaned dataset saved successfully as a single CSV file!")
```

Final cleaned dataset saved successfully as a single CSV file!

```
[18]: !gsutil mv gs://nypdbucket/notebooks/jupyter/cleaned_dataset/part-00000*.csv gs:
↪//nypdbucket/notebooks/jupyter/cleaned_dataset/final_cleaned_dataset.csv
```

Copying gs://nypdbucket/notebooks/jupyter/cleaned_dataset/part-00000-c5d2faff-641c-4b05-b267-072eb7c90d24-c000.csv [Content-Type=application/octet-stream]...
 Removing gs://nypdbucket/notebooks/jupyter/cleaned_dataset/part-00000-c5d2faff-641c-4b05-b267-072eb7c90d24-c000.csv...

Operation completed over 1 objects/372.5 MiB.

```
[19]: #https://storage.cloud.google.com/nypdbucket/notebooks/jupyter/nypd_dataset.csv
csv_file_path = "gs://nypdbucket/notebooks/jupyter/cleaned_dataset/
↪final_cleaned_dataset.csv"

# Read the CSV file into a DataFrame
df = (spark.read.format("csv")
      .option("header", "true") # If the CSV has a header row
      .option("inferSchema", "true") # Automatically infer the schema
      .load(csv_file_path))

# Show the first few rows of the DataFrame
df.show(5)
```

```
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+
|CREATE_DATE|INCIDENT_DATE|      INCIDENT_TIME|NYPD_PCT_CD|
BORO_NM|GEO_CD_X|GEO_CD_Y|RADIO_CODE|      TYP_DESC|CIP_JOBS|
ADD_TS|      DISP_TS|      ARRIVD_TS|      CLOSNG_TS| Latitude|
Longitude|dispatch_time|arrival_time|total_response_time|
```

```

+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
--+-----+-----+
| 01/01/2024| 12/31/2023|2024-12-17 22:31:06| 108| QUEENS| 1007298|
209993| 53S|VEHICLE ACCIDENT:...| Non CIP|2024-01-01 00:01:21|2024-01-01
00:02:19|2024-01-01 01:19:58|2024-01-01 01:20:02|40.743037|-73.916826|
0.97| 77.65| 78.62|
| 01/01/2024| 12/31/2023|2024-12-17 23:53:22| 114| QUEENS| 1002279|
222019| 34K1|ASSAULT (IN PROGR...|Critical|2024-01-01 00:06:11|2024-01-01
00:07:19|2024-01-01 00:19:27|2024-01-01 01:03:22|40.776057|-73.934906|
1.13| 12.13| 13.27|
| 01/01/2024| 12/31/2023|2024-12-17 23:59:17| 49| BRONX| 1020929|
254201| 11C4|ALARMS: COMMERCIA...| Non CIP|2024-01-01 00:04:51|2024-01-01
00:09:21|2024-01-01 00:15:11|2024-01-01 00:56:56| 40.86433|-73.867393|
4.5| 5.83| 10.33|
| 01/01/2024| 12/31/2023|2024-12-17 23:59:30| 34|MANHATTAN| 1003734|
253432| 11C4|ALARMS: COMMERCIA...| Non CIP|2024-01-01 00:04:57|2024-01-01
00:12:08|2024-01-01 00:29:16|2024-01-01 00:29:53|40.862274|-73.929562|
7.18| 17.13| 24.32|
| 01/01/2024| 01/01/2024|2024-12-17 00:00:14| 19|MANHATTAN| 992074|
217827| 53D|VEHICLE ACCIDENT:...| Non CIP|2024-01-01 00:00:14|2024-01-01
00:08:24|2024-01-01 00:36:32|2024-01-01 00:48:57|40.764566|-73.971757|
8.17| 28.13| 36.3|

```

```

+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
--+-----+-----+

```

only showing top 5 rows

[]:

```

[20]: counts = df_cleaned.select(
    [(F.sum(F.col(c).isNull().cast("int")).alias(c)) for c in df_cleaned.
     ↪columns]
)

# Show the count of null values for each column
counts.show()

```

[Stage 54:> (0 + 1) / 1]

```

+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+
|CREATE_DATE|INCIDENT_DATE|INCIDENT_TIME|NYPD_PCT_CD|BORO_NM|GEO_CD_X|GEO_CD_Y|R
ADIO_CODE|TYP_DESC|CIP_JOBS|ADD_TS|DISP_TS|ARRIVD_TS|CLOSNG_TS|Latitude|Longitud
e|dispatch_time|arrival_time|total_response_time|

```

```

+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+
|          0|          0|          0|          0|          0|          0|          0|          0|
0|          0|          0|          0|          0|          0|          0|          0|          0|
0|          0|          0|
+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+

```

1.1 Temporal Analysis

```

[21]: # Modify the INCIDENT_TIME column to extract only the time
df_cleaned_eda = df_cleaned.withColumn("INCIDENT_TIME", F.
    ↪date_format("INCIDENT_TIME", "HH:mm:ss"))\
    .withColumn("INCIDENT_DATE", F.
    ↪to_date("INCIDENT_DATE", "MM/dd/yyyy"))

# Show the updated DataFrame
df_cleaned.show(1)

```

```

+-----+-----+-----+-----+-----+-----+-----+
---+-----+-----+-----+-----+-----+-----+
---+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+
|CREATE_DATE|INCIDENT_DATE|
INCIDENT_TIME|NYPD_PCT_CD|BORO_NM|GEO_CD_X|GEO_CD_Y|RADIO_CODE|
TYP_DESC|CIP_JOBS|          ADD_TS|          DISP_TS|          ARRIVD_TS|
CLOSNG_TS| Latitude| Longitude|dispatch_time|arrival_time|total_response_time|
+-----+-----+-----+-----+-----+-----+-----+
---+-----+-----+-----+-----+-----+-----+
---+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+
| 01/01/2024| 12/31/2023|2024-12-17 22:31:06|          108| QUEENS| 1007298|
209993|          53S|VEHICLE ACCIDENT:...| Non CIP|2024-01-01 00:01:21|2024-01-01
00:02:19|2024-01-01 01:19:58|2024-01-01 01:20:02|40.743037|-73.916826|
0.97|          77.65|          78.62|
+-----+-----+-----+-----+-----+-----+-----+
---+-----+-----+-----+-----+-----+-----+
---+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+
only showing top 1 row

```

```
[22]: # Extract time-related features
df_cleaned_eda = df_cleaned_eda.withColumn("hour", F.hour("INCIDENT_TIME")) \
    .withColumn("day_of_week", F.dayofweek("INCIDENT_DATE")) \
    .withColumn("month", F.month("INCIDENT_DATE"))

# Hourly analysis
df_cleaned_eda.groupBy("hour").count().orderBy("hour").show()

# Day of the week analysis
df_cleaned_eda.groupBy("day_of_week").count().orderBy("day_of_week").show()

# Monthly analysis
df_cleaned_eda.groupBy("month").count().orderBy("month").show()
```

```
+----+-----+
|hour|count|
+----+-----+
|  0|66644|
|  1|59447|
|  2|51055|
|  3|43055|
|  4|40225|
|  5|36937|
|  6|39313|
|  7|46051|
|  8|72157|
|  9|73172|
| 10|71463|
| 11|71298|
| 12|71972|
| 13|71337|
| 14|68378|
| 15|62427|
| 16|81262|
| 17|78908|
| 18|78535|
| 19|76338|
+----+-----+
```

only showing top 20 rows

```
+-----+-----+
|day_of_week| count|
+-----+-----+
```


	1	212838
	2	220030
	3	215403
	4	218417
	5	214520
	6	220142
	7	217869
+-----+-----+		

```
[Stage 58:> (0 + 1) /
1]
```


month	count
	1 171869
	2 165935
	3 179985
	4 170371
	5 173827
	6 159930
	7 168275
	8 168419
	9 160604
	12 4
+-----+-----+	

```
[23]: # Convert to Pandas for visualization
day_of_week_data = df_cleaned_eda.groupby("day_of_week").count().toPandas()
month_data = df_cleaned_eda.groupby("month").count().toPandas()

# Plot day of week
day_of_week_data.plot(x="day_of_week", y="count", kind="bar", title="Incidents_
↳by Day of Week")
plt.show()
```



```
[24]: # Plot month
month_data.plot(x="month", y="count", kind="bar", title="Incidents by Month")
plt.show()
```

```
[25]: from pyspark.sql.functions import round

# Increase rounding precision for latitude and longitude
df_cleaned_eda = df_cleaned_eda.withColumn("latitude_rounded",
    ↪round("Latitude", 3)) \
    .withColumn("longitude_rounded", round("Longitude", 3))

# location_data = df_cleaned_eda.groupBy("latitude_rounded",
    ↪"longitude_rounded").count()
# location_data.orderBy("count", ascending=False).show(20, truncate=False)

# Select top 1000 locations with the highest incident counts
top_locations = df_cleaned_eda.groupBy("latitude_rounded", "longitude_rounded").
    ↪count().orderBy("count", ascending=False).limit(40000).toPandas()
```

```
[26]: import folium
from folium.plugins import HeatMap
```

```

# Convert the top_locations DataFrame into a format suitable for HeatMap
heatmap_data = [[row["latitude_rounded"], row["longitude_rounded"],
    ↳row["count"]] for index, row in top_locations.iterrows()]

# Create a base map centered at the mean latitude and longitude
m = folium.Map(
    location=[top_locations["latitude_rounded"].mean(),
    ↳top_locations["longitude_rounded"].mean()],
    zoom_start=12
)

# Add the heatmap layer
HeatMap(heatmap_data, radius=10, blur=15, max_zoom=1).add_to(m)

# Save and display the map
m.save("top_locations_heatmap.html")
m

```

[26]: <folium.folium.Map at 0x7fe0888ac250>

2 EDA FOR INCIDENT CLASSIFICATION

```

[27]: # Count incidents by type
incident_type_analysis = df_cleaned_eda.groupby("TYP_DESC").count().
    ↳orderBy("count", ascending=False)

# Display the top incident types
incident_type_analysis.show(10, truncate=False)

# Convert incident type analysis to Pandas for plotting
incident_type_pd = incident_type_analysis.limit(10).toPandas()

# Plot the bar chart
plt.figure(figsize=(10, 6))
plt.bar(incident_type_pd["TYP_DESC"], incident_type_pd["count"],
    ↳color="skyblue")
plt.title("Top 10 Incident Types")
plt.xlabel("Incident Type")
plt.ylabel("Count")
plt.xticks(rotation=45)
plt.show()

```

```


+-----+-----+
|TYP_DESC| count|
+-----+-----+

```

VISIBILITY PATROL: DIRECTED	85811
SEE COMPLAINANT: OTHER/INSIDE	81903
AMBULANCE CASE: EDP/INSIDE	69473
DISPUTE: FAMILY	62591
INVESTIGATE/POSSIBLE CRIME: CALLS FOR HELP/INSIDE	59328
INVESTIGATE/POSSIBLE CRIME: SERIOUS/OTHER	53826
ALARMS: COMMERCIAL/BURGLARY	53799
OTHER CRIMES (IN PROGRESS): HARASSMENT/INSIDE	50462
DISPUTE: INSIDE	48050
TRANSIT PATROL/INSPECTION BY NON-TRANSIT BUREAU PERSONNEL	39887

+-----+

only showing top 10 rows




```
[28]: # Count incidents by radio code
radio_code_analysis = df_cleaned_eda.groupBy("RADIO_CODE").count().
    ↪orderBy("count", ascending=False)

# Display the top radio codes
radio_code_analysis.show(10, truncate=False)

# Convert radio code analysis to Pandas for plotting
radio_code_pd = radio_code_analysis.limit(10).toPandas()

# Plot the bar chart
plt.figure(figsize=(10, 6))
plt.bar(radio_code_pd["RADIO_CODE"], radio_code_pd["count"], color="lightgreen")
plt.title("Top 10 Radio Codes")
plt.xlabel("Radio Code")
plt.ylabel("Count")
plt.xticks(rotation=45)
plt.show()
```

```
+-----+-----+
|RADIO_CODE|count|
+-----+-----+
|75D       |85811|
|68Q1      |81903|
|54E1      |69473|
|52D6      |62591|
|10H1      |59328|
|10Y3      |53826|
|11C4      |53799|
|39H1      |50462|
|52D1      |48050|
|75T       |39887|
+-----+-----+
only showing top 10 rows
```



```
[29]: # Count incidents by job classification
job_classification_analysis = df_cleaned_eda.groupby("CIP_JOBS").count().
    ↳orderBy("count", ascending=False)

# Display the top job classifications
job_classification_analysis.show(10, truncate=False)

# Convert job classification analysis to Pandas for plotting
job_classification_pd = job_classification_analysis.limit(10).toPandas()

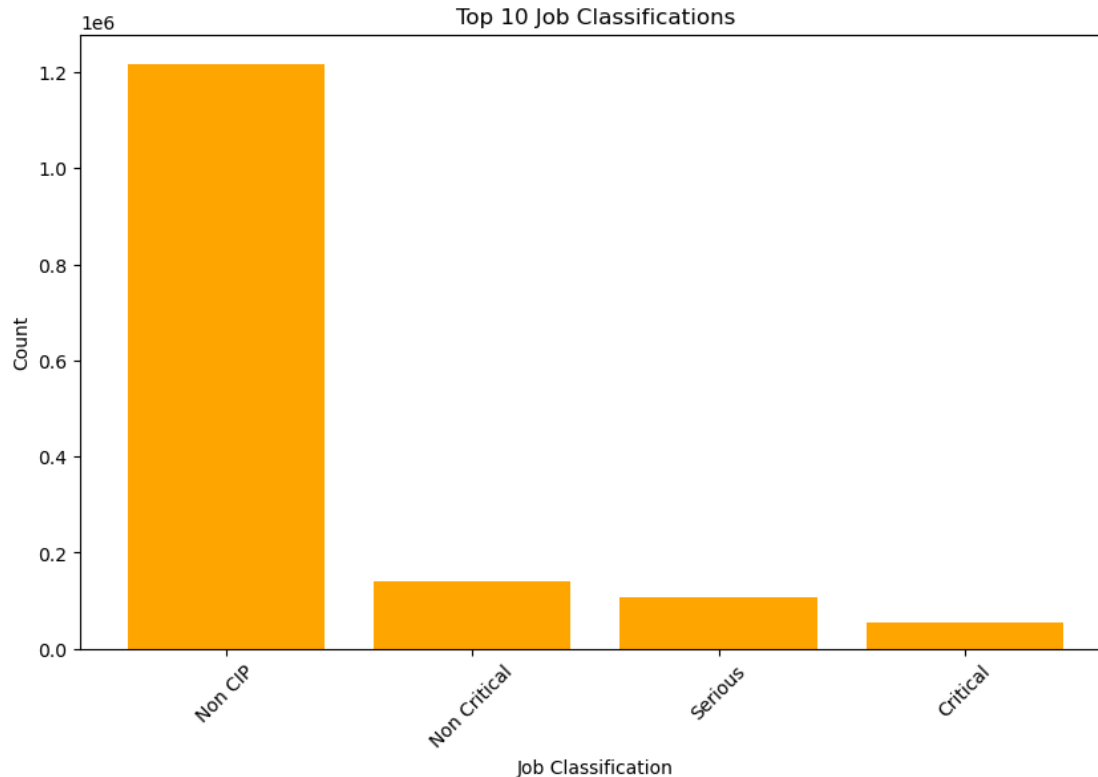
# Plot the bar chart
plt.figure(figsize=(10, 6))
plt.bar(job_classification_pd["CIP_JOBS"], job_classification_pd["count"],
    ↳color="orange")
plt.title("Top 10 Job Classifications")
plt.xlabel("Job Classification")
plt.ylabel("Count")
plt.xticks(rotation=45)
plt.show()
```

```
+-----+-----+
|CIP_JOBS |count |
```

```

+-----+-----+
|Non CIP      |1217682|
|Non Critical|140618 |
|Serious      |106119 |
|Critical     |54800  |
+-----+-----+

```



```

[30]: # Group by borough and incident type
incident_types_by_borough = df_cleaned_eda.groupBy("BORO_NM", "TYP_DESC").
    ↪count().orderBy("count", ascending=False)

# Display the top borough-incident type combinations
incident_types_by_borough.show(10, truncate=False)

# Convert to Pandas for plotting
incident_types_pd = incident_types_by_borough.limit(10).toPandas()

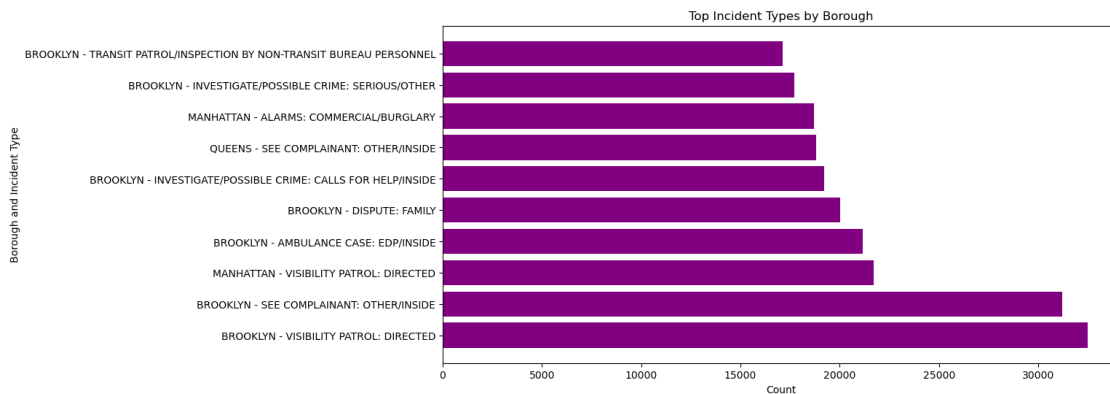
# Plot bar chart
plt.figure(figsize=(12, 6))

```

```
plt.barh(incident_types_pd["BORO_NM"] + " - " + incident_types_pd["TYP_DESC"],
         incident_types_pd["count"], color="purple")
plt.title("Top Incident Types by Borough")
plt.xlabel("Count")
plt.ylabel("Borough and Incident Type")
plt.show()
```

```
+-----+-----+-----+-----+
|BORO_NM |TYP_DESC                                     |count|
+-----+-----+-----+-----+
|BROOKLYN |VISIBILITY PATROL: DIRECTED                     |32476|
|BROOKLYN |SEE COMPLAINANT: OTHER/INSIDE                   |31219|
|MANHATTAN|VISIBILITY PATROL: DIRECTED                     |21704|
|BROOKLYN |AMBULANCE CASE: EDP/INSIDE                     |21156|
|BROOKLYN |DISPUTE: FAMILY                                |20031|
|BROOKLYN |INVESTIGATE/POSSIBLE CRIME: CALLS FOR HELP/INSIDE |19206|
|QUEENS   |SEE COMPLAINANT: OTHER/INSIDE                   |18806|
|MANHATTAN|ALARMS: COMMERCIAL/BURGLARY                     |18684|
|BROOKLYN |INVESTIGATE/POSSIBLE CRIME: SERIOUS/OTHER       |17696|
|BROOKLYN |TRANSIT PATROL/INSPECTION BY NON-TRANSIT BUREAU PERSONNEL |17104|
+-----+-----+-----+-----+
```

only showing top 10 rows



```
[31]: # Average response times by borough
response_time_by_borough = df_cleaned_eda.groupby("BORO_NM").
    .mean("dispatch_time", "arrival_time", "total_response_time")
response_time_by_borough.show(truncate=False)
```

[Stage 70:> (0 + 1) / 1]

BORO_NM	avg(dispatch_time)	avg(arrival_time)	avg(total_response_time)
QUEENS	3.4332315183604516	19.172429105356045	22.60586888162176
BRONX	3.742271012683865	17.651551254424962	21.393927382916427
MANHATTAN	3.0732346362753376	15.06390266740991	18.1371637638107
BROOKLYN	2.9635499758643937	16.63173771245205	19.595526404670345
STATEN ISLAND	2.848513425194413	15.632137955451899	18.480836147269635
(null)	1.568095238095238	13.253333333333334	14.820952380952384

```
[32]: from pyspark.sql.functions import hour, dayofweek
```

```
# Add hour and day of the week columns
df_cleaned_eda = df_cleaned_eda.withColumn("hour", hour("INCIDENT_TIME")) \
    .withColumn("day_of_week",
        dayofweek("INCIDENT_DATE"))

# Average response times by hour
response_time_by_hour = df_cleaned_eda.groupBy("hour").mean("dispatch_time",
    "arrival_time", "total_response_time")
response_time_by_hour.show(truncate=False)

# Average response times by day of the week
response_time_by_day = df_cleaned_eda.groupBy("day_of_week").
    mean("dispatch_time", "arrival_time", "total_response_time")
response_time_by_day.show(truncate=False)
```

hour	avg(dispatch_time)	avg(arrival_time)	avg(total_response_time)
22	3.1272606028272865	13.53621373542937	16.66363704831049
23	4.341739894208651	18.10833291043063	22.450261354985987
0	2.7815602304781795	14.271719284559152	17.05347188043951
1	2.467993675038265	12.777316433127371	15.245506922132085
2	2.3968036431300566	12.546193712664378	14.94316325531275
3	2.4486705376843902	12.623134130764058	15.07200092904415
4	2.4914846488502476	13.001222871349647	15.492924549409468
5	2.507468121395905	13.641176868723159	16.148846684895787
6	2.6962617963524047	14.082330018061256	16.77873069976859
7	4.150677509717355	23.767868667348697	27.918648237823497
8	2.567756004268294	18.62475393932958	21.192663220479318
9	2.860873968184314	18.695965533266726	21.55695238616037
10	3.125105299245533	18.78806711165481	21.913313322979636

11	3.2958690285840233	19.24268983702451	22.538693511740846	
12	3.4051720113375272	19.536395959542755	22.94173150669866	
13	3.3526879459464434	18.077859876364244	21.43069893603734	
14	3.3539938284241235	15.20779753722193	18.561966129457435	
15	4.8294797123040105	20.968853380750698	25.798443622151325	
16	3.327067386970393	17.853575964168037	21.180765302356836	
17	3.28389922441303	17.46224476605933	20.746284533888502	

only showing top 20 rows

[Stage 72:> (0 + 1) / 1]

day_of_week	avg(dispatch_time)	avg(arrival_time)	avg(total_response_time)
1	3.2058363638085847	16.66860969375259	19.87459903776274
2	3.2196181884301733	17.178221288001104	20.39801754306018
3	3.1854086990445323	16.848214973788462	20.03376726414887
4	3.1067741064125314	16.789266403250615	19.89619612941861
5	3.23745730002018	16.945285148232344	20.18288635091962
6	3.3220365491380233	17.247511106462397	20.569702146793453
7	3.2316336881350463	16.921222294125407	20.15301777673505

```
[33]: # Average response times by incident type
response_time_by_incident = df_cleaned_eda.groupby("TYP_DESC").
    ↪mean("dispatch_time", "arrival_time", "total_response_time")
response_time_by_incident.orderBy("avg(total_response_time)", ascending=False).
    ↪show(10, truncate=False)
```

[Stage 73:> (0 + 1) / 1]

TYP_DESC	avg(dispatch_time)	avg(arrival_time)	avg(total_response_time)
CELLULAR OPEN LINE		14.65	60.17
74.82			
LARCENY (PAST): VEHICLE/SCHOOL		7.983333333333333	
50.50666666666666	58.48666666666666		
DISORDERLY: NOISE/OUTSIDE		8.43	46.18
54.62			

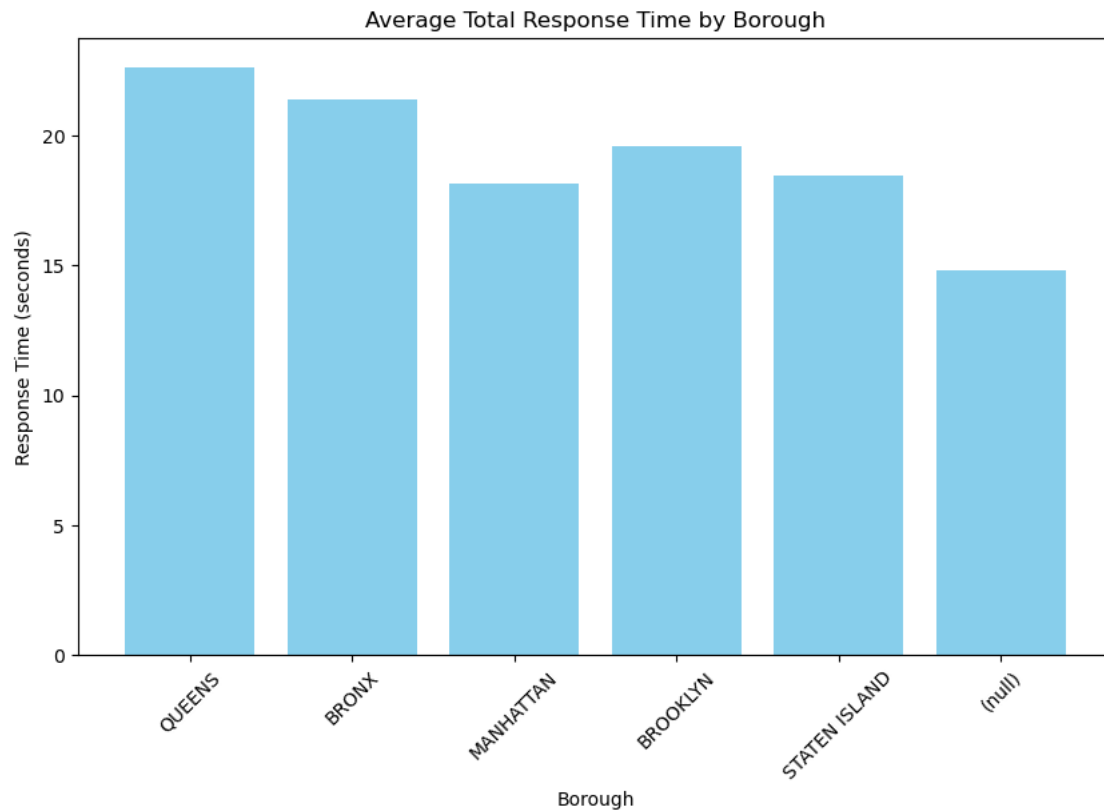
UTILITY TROUBLE (SPECIFY): LTD ACC HWY	1.2	52.85
54.05		
ASSAULT (PAST): OTHER/LTD ACC HWY	8.233333333333333	
41.54333333333333 49.776666666666664		
LARCENY (PAST): OTHER/LTD ACC HWY	8.540000000000001	39.725
48.26500000000001		
SUSP PACKAGE: LTD ACC HWY	6.823333333333334	
41.22166666666667 48.04500000000001		
INVESTIGATE/POSSIBLE CRIME: NARCO SALES/LTD ACC HWY	2.58	45.2
47.78		
DISABLED VEHICLE: LTD ACC HWY	8.128235294117646	
38.980000000000004 47.107058823529414		
INVESTIGATE/POSSIBLE CRIME: SUSP VEHICLE/TRANSIT	4.0525	
42.620000000000005 46.675		

only showing top 10 rows

```
[34]: import matplotlib.pyplot as plt


# Convert to Pandas for plotting
response_time_borough_pd = response_time_by_borough.toPandas()

# Plot bar chart
plt.figure(figsize=(10, 6))
plt.bar(response_time_borough_pd["BORO_NM"],
        response_time_borough_pd["avg(total_response_time)"], color="skyblue")
plt.title("Average Total Response Time by Borough")
plt.xlabel("Borough")
plt.ylabel("Response Time (seconds)")
plt.xticks(rotation=45)
plt.show()
```



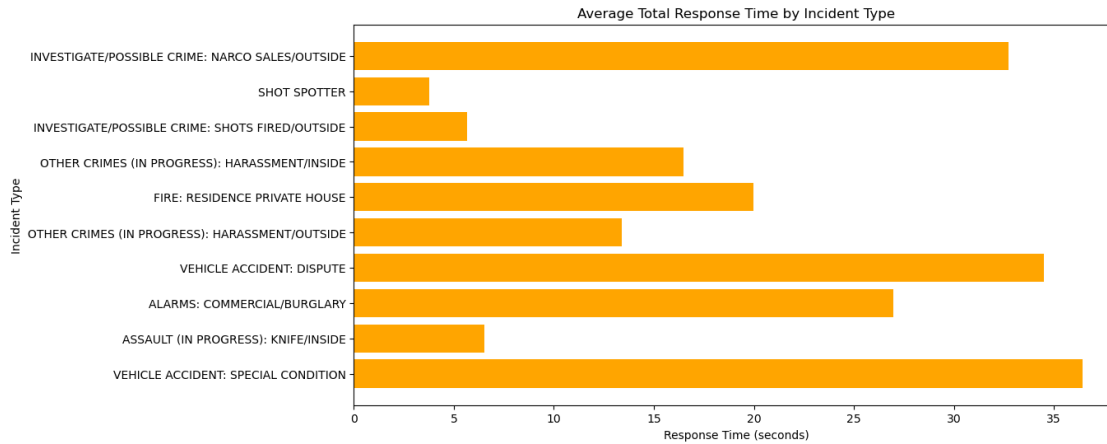
```
[35]: # Convert to Pandas for plotting
response_time_hour_pd = response_time_by_hour.toPandas()

# Plot line chart
plt.figure(figsize=(10, 6))
plt.plot(response_time_hour_pd["hour"], □
         ↪ response_time_hour_pd["avg(total_response_time)"], marker="o", color="green")
plt.title("Average Total Response Time by Hour of Day")
plt.xlabel("Hour of Day")
plt.ylabel("Response Time (seconds)")
plt.grid()
plt.show()
```



```
[36]: # Convert to Pandas for plotting
response_time_incident_pd = response_time_by_incident.limit(10).toPandas()

# Plot horizontal bar chart
plt.figure(figsize=(12, 6))
plt.barh(response_time_incident_pd["TYP_DESC"],
          response_time_incident_pd["avg(total_response_time)"], color="orange")
plt.title("Average Total Response Time by Incident Type")
plt.xlabel("Response Time (seconds)")
plt.ylabel("Incident Type")
plt.show()
```



3 Modeling

3.1 Linear Regression

```
[37]: from pyspark.ml.feature import VectorAssembler
      from pyspark.ml.regression import LinearRegression
      from pyspark.ml.evaluation import RegressionEvaluator
      from pyspark.sql import SparkSession
      from pyspark.ml.regression import GBRegressor
```

```
[38]: # Feature columns for the regression model
      feature_columns = ["dispatch_time", "arrival_time"]

      # Assemble features into a single vector
      assembler = VectorAssembler(inputCols=feature_columns, outputCol="features")

      # Transform the data
      df_prepared = assembler.transform(df_cleaned).select("features",
      ↪ "total_response_time")

      # Show the prepared data
      print("Prepared Data for Linear Regression:")
      df_prepared.show(5, truncate=False)
```

Prepared Data for Linear Regression:

features	total_response_time
[0.97,77.65]	78.62
[1.13,12.13]	13.27
[4.5,5.83]	10.33
[7.18,17.13]	24.32

```
| [8.17, 28.13] | 36.3 |
+-----+-----+
only showing top 5 rows
```

```
[39]: # Split the data into train and test sets (80% train, 20% test)
train_data, test_data = df_prepared.randomSplit([0.6, 0.4], seed=42)

print("Training Data Count:", train_data.count())
print("Test Data Count:", test_data.count())
```

Training Data Count: 912778

```
[Stage 79:> (0 + 1) / 1]
```

Test Data Count: 606441

```
[40]: # Initialize the Linear Regression model
lr = LinearRegression(featuresCol="features", labelCol="total_response_time",
    ↪ predictionCol="prediction")

# Train the model
lr_model = lr.fit(train_data)

# Print model coefficients and intercept
print("Coefficients:", lr_model.coefficients)
print("Intercept:", lr_model.intercept)
```

24/12/17 01:37:02 WARN Instrumentation: [901bbe0c] regParam is zero, which might cause numerical instability and overfitting.

```
[Stage 81:> (0 + 1) / 1]
```

Coefficients: [0.9999813888156942, 0.9999967289283758]

Intercept: 0.00026806356868136423

```
[41]: # Predict on test data
predictions = lr_model.transform(test_data)

# Show predictions
print("Linear Regression Predictions:")
predictions.select("features", "total_response_time", "prediction").show(10,
    ↪ truncate=False)

# Evaluate the model
```

```

evaluator = RegressionEvaluator(labelCol="total_response_time",
    predictionCol="prediction", metricName="rmse")

# Calculate RMSE
rmse = evaluator.evaluate(predictions)
print("Root Mean Squared Error (RMSE):", rmse)

# Optionally calculate R2
evaluator_r2 = RegressionEvaluator(labelCol="total_response_time",
    predictionCol="prediction", metricName="r2")
r2 = evaluator_r2.evaluate(predictions)
print("R2 Score:", r2)

```

Linear Regression Predictions:

```

+-----+-----+-----+
|features |total_response_time|prediction      |
+-----+-----+-----+
| [0.02,0.02] |0.03              |0.04026762592356276|
| [0.02,0.02] |0.03              |0.04026762592356276|
| [0.02,0.02] |0.03              |0.04026762592356276|
| [0.02,0.02] |0.03              |0.04026762592356276|
| [0.02,0.02] |0.03              |0.04026762592356276|
| [0.02,0.02] |0.03              |0.04026762592356276|
| [0.02,0.02] |0.03              |0.04026762592356276|
| [0.02,0.02] |0.03              |0.04026762592356276|
| [0.02,0.02] |0.03              |0.04026762592356276|
| [0.02,0.02] |0.03              |0.04026762592356276|
+-----+-----+-----+
only showing top 10 rows

```

Root Mean Squared Error (RMSE): 0.004880453984611745

```

[Stage 84:>                                     (0 + 1) /
1]
R2 Score: 0.9999999572697764

```

```

[42]: # Model summary
training_summary = lr_model.summary

print("Training Summary:")
print("RMSE on Training Data:", training_summary.rootMeanSquaredError)
print("R2 on Training Data:", training_summary.r2)

```

Training Summary:
RMSE on Training Data: 0.00488897638958177
R2 on Training Data: 0.9999999571389353

3.2 Random Forest

```
[43]: from pyspark.ml.regression import RandomForestRegressor

# Initialize Random Forest Regressor
rf = RandomForestRegressor(featuresCol="features",
    ↳labelCol="total_response_time", predictionCol="prediction", numTrees=100)

# Train the model
rf_model = rf.fit(train_data)

print("Random Forest model training completed!")
```

Random Forest model training completed!

```
[44]: # Predict on the test data
predictions = rf_model.transform(test_data)

# Show predictions
print("Predictions on Test Data:")
predictions.select("features", "total_response_time", "prediction").show(10,
    ↳truncate=False)
```

Predictions on Test Data:

WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.spark.util.SizeEstimator\$
(file:/usr/lib/spark/jars/spark-core_2.12-3.5.1.jar) to field
java.nio.charset.Charset.name
WARNING: Please consider reporting this to the maintainers of
org.apache.spark.util.SizeEstimator\$
WARNING: Use --illegal-access=warn to enable warnings of further illegal
reflective access operations
WARNING: All illegal access operations will be denied in a future release
[Stage 99:> (0 + 1) / 1]

```
+-----+-----+-----+
|features|total_response_time|prediction|
+-----+-----+-----+
|[0.02,0.02]|0.03|0.7888153334353398|
|[0.02,0.02]|0.03|0.7888153334353398|
|[0.02,0.02]|0.03|0.7888153334353398|
|[0.02,0.02]|0.03|0.7888153334353398|
|[0.02,0.02]|0.03|0.7888153334353398|
```


[0.02,0.02] 0.03	0.7888153334353398
[0.02,0.02] 0.03	0.7888153334353398
[0.02,0.02] 0.03	0.7888153334353398
[0.02,0.02] 0.03	0.7888153334353398
[0.02,0.02] 0.03	0.7888153334353398
+-----+	+-----+

only showing top 10 rows

```
[45]: # Initialize evaluators
evaluator_rmse = RegressionEvaluator(labelCol="total_response_time",
    predictionCol="prediction", metricName="rmse")
evaluator_r2 = RegressionEvaluator(labelCol="total_response_time",
    predictionCol="prediction", metricName="r2")

# Calculate RMSE and R2
rmse = evaluator_rmse.evaluate(predictions)
r2 = evaluator_r2.evaluate(predictions)

print(f"Root Mean Squared Error (RMSE): {rmse}")
print(f"R2 Score: {r2}")
```

[Stage 101:> (0 + 1) / 1]

Root Mean Squared Error (RMSE): 3.678746451618793
R2 Score: 0.9757219146809109

3.3 Gradient Boosting

```
[46]: # Initialize Gradient Boosted Tree Regressor
gbt = GBRegressor(featuresCol="features", labelCol="total_response_time",
    predictionCol="prediction", maxIter=100)

# Train the model
gbt_model = gbt.fit(train_data)

print("Gradient Boosting model training completed!")
```

Gradient Boosting model training completed!

```
[47]: # Predict on the test data
predictions = gbt_model.transform(test_data)

# Show predictions
```

```
print("Predictions on Test Data:")
predictions.select("features", "total_response_time", "prediction").show(10,
↳truncate=False)
```

Predictions on Test Data:

[Stage 1106:> (0 + 1) / 1]

features	total_response_time	prediction
[0.02,0.02]	0.03	0.032046177496824646
[0.02,0.02]	0.03	0.032046177496824646
[0.02,0.02]	0.03	0.032046177496824646
[0.02,0.02]	0.03	0.032046177496824646
[0.02,0.02]	0.03	0.032046177496824646
[0.02,0.02]	0.03	0.032046177496824646
[0.02,0.02]	0.03	0.032046177496824646
[0.02,0.02]	0.03	0.032046177496824646
[0.02,0.02]	0.03	0.032046177496824646
[0.02,0.02]	0.03	0.032046177496824646
[0.02,0.02]	0.03	0.032046177496824646

only showing top 10 rows

```
[48]: # Initialize evaluators
evaluator_rmse = RegressionEvaluator(labelCol="total_response_time",
↳predictionCol="prediction", metricName="rmse")
evaluator_r2 = RegressionEvaluator(labelCol="total_response_time",
↳predictionCol="prediction", metricName="r2")

# Calculate RMSE and R2
rmse = evaluator_rmse.evaluate(predictions)
r2 = evaluator_r2.evaluate(predictions)

print(f"Root Mean Squared Error (RMSE): {rmse}")
print(f"R2 Score: {r2}")
```

[Stage 1108:> (0 + 1) / 1]

Root Mean Squared Error (RMSE): 1.5631399430441817
R2 Score: 0.9956166066697358

```
[49]: from pyspark.sql.functions import when, col
from pyspark.ml.evaluation import MulticlassClassificationEvaluator
```

```

from pyspark.ml.classification import RandomForestClassifier
from pyspark.ml.feature import VectorAssembler

# Step 1: Categorize total_response_time into binary classes
threshold = 60 # Set a threshold for classification
df_classification = df_cleaned.withColumn("label",
    when(col("total_response_time") > threshold, 1).otherwise(0))

# Assemble features
feature_columns = ["dispatch_time", "arrival_time"] # Include relevant features
assembler = VectorAssembler(inputCols=feature_columns, outputCol="features")
df_prepared = assembler.transform(df_classification).select("features", "label")

# Split into train and test sets
train_data, test_data = df_prepared.randomSplit([0.6, 0.4], seed=42)

```

```

[50]: # Initialize Random Forest Classifier
rf_classifier = RandomForestClassifier(featuresCol="features",
    labelCol="label", predictionCol="prediction", numTrees=100)

# Train the classifier
rf_model = rf_classifier.fit(train_data)

# Make predictions on test data
predictions = rf_model.transform(test_data)

# Show predictions
predictions.select("features", "label", "prediction").show(10)

```

[Stage 1124:>

(0 + 1) / 1]

```

+-----+-----+-----+
| features|label|prediction|
+-----+-----+-----+
|[0.02,0.02]| 0| 0.0|
|[0.02,0.02]| 0| 0.0|
|[0.02,0.02]| 0| 0.0|
|[0.02,0.02]| 0| 0.0|
|[0.02,0.02]| 0| 0.0|
|[0.02,0.02]| 0| 0.0|
|[0.02,0.02]| 0| 0.0|
|[0.02,0.02]| 0| 0.0|
|[0.02,0.02]| 0| 0.0|
|[0.02,0.02]| 0| 0.0|
+-----+-----+-----+

```

only showing top 10 rows

```
[51]: # Accuracy
evaluator_accuracy = MulticlassClassificationEvaluator(labelCol="label",
    ↪predictionCol="prediction", metricName="accuracy")
accuracy = evaluator_accuracy.evaluate(predictions)

# Precision
evaluator_precision = MulticlassClassificationEvaluator(labelCol="label",
    ↪predictionCol="prediction", metricName="weightedPrecision")
precision = evaluator_precision.evaluate(predictions)

# Recall
evaluator_recall = MulticlassClassificationEvaluator(labelCol="label",
    ↪predictionCol="prediction", metricName="weightedRecall")
recall = evaluator_recall.evaluate(predictions)

# Print the metrics
print(f"Accuracy: {accuracy}")
print(f"Precision: {precision}")
print(f"Recall: {recall}")
```

[Stage 1129:>

(0 + 1) / 1]

Accuracy: 0.9904821738635745

Precision: 0.9903975325539243

Recall: 0.9904821738635745

3.4 K-Means clustering

```
[52]: from pyspark.ml.clustering import KMeans
from pyspark.ml.feature import VectorAssembler
from pyspark.sql.functions import col
```

```
[53]: # Select relevant columns (latitude and longitude)
selected_features = ["Latitude", "Longitude"] # Add other relevant features if
    ↪needed

# Assemble features into a single vector
assembler = VectorAssembler(inputCols=selected_features, outputCol="features")

# Transform the data
df_prepared = assembler.transform(df_cleaned).select("features",
    ↪*selected_features)

# Show prepared data
print("Prepared Data for K-means Clustering:")
df_prepared.show(5, truncate=False)
```

Prepared Data for K-means Clustering:

```
+-----+-----+-----+
|features          |Latitude |Longitude |
+-----+-----+-----+
|[40.743037,-73.916826]|40.743037|-73.916826|
|[40.776057,-73.934906]|40.776057|-73.934906|
|[40.86433,-73.867393]|40.86433 |-73.867393|
|[40.862274,-73.929562]|40.862274|-73.929562|
|[40.764566,-73.971757]|40.764566|-73.971757|
+-----+-----+-----+
```

only showing top 5 rows

```
[54]: # Initialize K-means model
k = 5 # Define the number of clusters
kmeans = KMeans(featuresCol="features", predictionCol="cluster", k=k, seed=42)

# Train the K-means model
kmeans_model = kmeans.fit(df_prepared)

# Assign clusters to data points
df_clusters = kmeans_model.transform(df_prepared)

# Show the resulting clusters
print("Data with Cluster Assignments:")
df_clusters.select("Latitude", "Longitude", "cluster").show(10)
```

Data with Cluster Assignments:

```
+-----+-----+-----+
| Latitude| Longitude|cluster|
+-----+-----+-----+
|40.743037|-73.916826|4|
|40.776057|-73.934906|4|
|40.86433|-73.867393|0|
|40.862274|-73.929562|0|
|40.764566|-73.971757|4|
|40.706102|-73.793242|2|
|40.740547|-74.008547|4|
|40.706528|-73.791997|2|
|40.770813|-73.811147|2|
|40.70215|-73.790564|2|
+-----+-----+-----+
```

only showing top 10 rows

```
[55]: # Get cluster centers
centers = kmeans_model.clusterCenters()
print("Cluster Centers (High-Risk Zones):")
for idx, center in enumerate(centers):
    print(f"Cluster {idx}: {center}")
```


```
Cluster Centers (High-Risk Zones):
Cluster 0: [ 40.83743223 -73.90282894]
Cluster 1: [ 40.60199467 -74.12495999]
Cluster 2: [ 40.70022946 -73.81260808]
Cluster 3: [ 40.65558713 -73.95103755]
Cluster 4: [ 40.74828972 -73.96746351]
```

```
[56]: import matplotlib.pyplot as plt

# Convert to Pandas for plotting
df_pandas = df_clusters.select("Latitude", "Longitude", "cluster").toPandas()

# Plot clusters
plt.figure(figsize=(10, 6))
for cluster_id in df_pandas["cluster"].unique():
    cluster_data = df_pandas[df_pandas["cluster"] == cluster_id]
    plt.scatter(cluster_data["Longitude"], cluster_data["Latitude"],
        label=f"Cluster {cluster_id}")

plt.title("K-means Clustering of High-Risk Zones")
plt.xlabel("Longitude")
plt.ylabel("Latitude")
plt.legend()
plt.show()
```



```
[57]: from pyspark.sql.functions import col, sqrt, pow, mean
      from pyspark.ml.linalg import Vectors
      import numpy as np
```

```
[58]: from pyspark.sql.functions import col, udf
      from pyspark.sql.types import DoubleType
      from pyspark.ml.linalg import DenseVector
      import math
      from pyspark import SparkContext

      # Step 2.1: Get cluster centers
      centers = kmeans_model.clusterCenters()
      cluster_centers = {i: centers[i] for i in range(len(centers))}

      # Step 2.2: Broadcast cluster centers for efficiency
      sc = SparkContext.getOrCreate()
      broadcast_centers = sc.broadcast(cluster_centers)

      # Step 2.3: Define a UDF to calculate Euclidean distance
      def euclidean_distance(point, cluster_id):
          center = broadcast_centers.value[cluster_id] # Fetch cluster center from
          ↪ broadcasted dictionary
          return float(math.sqrt(sum((point[i] - center[i]) ** 2 for i in
          ↪ range(len(center)))))
```

```

# Register the UDF
distance_udf = udf(euclidean_distance, DoubleType())

# Step 2.4: Add intra-cluster distances to the DataFrame
df_with_distance = df_clusters.withColumn(
    "intra_cluster_distance",
    distance_udf(col("features"), col("cluster"))
)

# Show the updated DataFrame
print("Data with Intra-Cluster Distances:")
df_with_distance.select("features", "cluster", "intra_cluster_distance").
    ↪show(10, truncate=False)

```

Data with Intra-Cluster Distances:

```
[Stage 1184:>                                                                    (0 + 1) /
1]
```

features	cluster	intra_cluster_distance
[40.743037,-73.916826]	4	0.050909216914502266
[40.776057,-73.934906]	4	0.04279034383589714
[40.86433,-73.867393]	0	0.04448815501335516
[40.862274,-73.929562]	0	0.036493420205705035
[40.764566,-73.971757]	4	0.016833048236449075
[40.706102,-73.793242]	2	0.02023688849743139
[40.740547,-74.008547]	4	0.041806732023081376
[40.706528,-73.791997]	2	0.0215519858106137
[40.770813,-73.811147]	2	0.07059866101884622
[40.70215,-73.790564]	2	0.022127579853446676

only showing top 10 rows

```

[59]: import numpy as np

# Function to compute pairwise distances between cluster centers
def pairwise_distances(centers):
    k = len(centers)
    distances = np.zeros((k, k))
    for i in range(k):
        for j in range(k):
            if i != j:

```



```

        distances[i][j] = np.linalg.norm(np.array(centers[i]) - np.
↪array(centers[j]))
        return distances

# Compute pairwise distances
inter_cluster_distances = pairwise_distances(centers)

# Print pairwise inter-cluster distances
print("Inter-Cluster Distances:")
print(inter_cluster_distances)

```

Inter-Cluster Distances:

```

[[0.          0.32368665 0.16420842 0.18812685 0.11010911]
 [0.32368665 0.          0.32743517 0.18199222 0.21495902]
 [0.16420842 0.32743517 0.          0.14544984 0.16214189]
 [0.18812685 0.18199222 0.14544984 0.          0.09414659]
 [0.11010911 0.21495902 0.16214189 0.09414659 0.          ]]

```

```

[60]: # Collect intra-cluster dispersions
dispersion_values = df_with_distance.groupby("cluster").
↪agg(mean("intra_cluster_distance").alias("dispersion")).collect()

# Convert dispersion values into a dictionary
dispersion_dict = {row['cluster']: row['dispersion'] for row in_
↪dispersion_values}

# Compute Davies-Bouldin Index
db_index = 0
k = len(centers)

for i in range(k):
    max_ratio = 0
    for j in range(k):
        if i != j:
            ratio = (dispersion_dict[i] + dispersion_dict[j]) /_
↪inter_cluster_distances[i][j]
            max_ratio = max(max_ratio, ratio)
    db_index += max_ratio

db_index /= k

print(f"Davies-Bouldin Index (DBI): {db_index}")

```

[Stage 1185:>

(0 + 1) /

1]

Davies-Bouldin Index (DBI): 0.7477194917382969

```
[61]: from pyspark.ml.evaluation import ClusteringEvaluator
from pyspark.ml.clustering import KMeans
from pyspark.ml.feature import VectorAssembler

# Step 1: Assemble Features (if not already done)
feature_columns = ["Latitude", "Longitude"] # Include spatial or relevant
# features
assembler = VectorAssembler(inputCols=feature_columns, outputCol="features")
df_prepared = assembler.transform(df_cleaned)

# Step 2: Train K-means Model
k = 5 # Number of clusters
kmeans = KMeans(featuresCol="features", predictionCol="prediction", k=k,
# seed=42)
kmeans_model = kmeans.fit(df_prepared)

# Step 3: Assign Clusters to Data Points
df_clusters = kmeans_model.transform(df_prepared)

# Step 4: Compute Silhouette Score
evaluator = ClusteringEvaluator(featuresCol="features",
# predictionCol="prediction", metricName="silhouette",
# distanceMeasure="squaredEuclidean")

silhouette_score = evaluator.evaluate(df_clusters)

print(f"Silhouette Score: {silhouette_score}")
```

[Stage 1240:>

(0 + 1) / 1]

Silhouette Score: 0.5913598921815876

4 Scale in Scale out

```
[62]: import time
from pyspark.ml.regression import RandomForestRegressor, GBTRRegressor
from pyspark.ml.feature import VectorAssembler
from pyspark.ml.evaluation import RegressionEvaluator

# Create data subsets at 20%, 50%, 75%, and 100% of the full dataset
fractions = {"20%": 0.2, "50%": 0.5, "75%": 0.75, "100%": 1.0}
data_subsets = {}

# Generate subsets and store them in a dictionary
for key, fraction in fractions.items():
    data_subsets[key] = df_cleaned.sample(fraction=fraction, seed=42)
```

```
print("Data subsets created for 20%, 50%, 75%, and 100%")
```

Data subsets created for 20%, 50%, 75%, and 100%

```
[63]: # Initialize models
rf = RandomForestRegressor(featuresCol="features",
    ↳ labelCol="total_response_time", predictionCol="prediction", numTrees=100)
gbt = GBTRegressor(featuresCol="features", labelCol="total_response_time",
    ↳ predictionCol="prediction", maxIter=100)

# Initialize evaluator
evaluator_rmse = RegressionEvaluator(labelCol="total_response_time",
    ↳ predictionCol="prediction", metricName="rmse")
evaluator_r2 = RegressionEvaluator(labelCol="total_response_time",
    ↳ predictionCol="prediction", metricName="r2")

# Results storage
results = []

# Train and evaluate both models on each subset
for size, df_subset in data_subsets.items():
    # Assemble features
    assembler = VectorAssembler(inputCols=["dispatch_time", "arrival_time"],
    ↳ outputCol="features")
    df_prepared = assembler.transform(df_subset).select("features",
    ↳ "total_response_time")

    # Split data
    train_data, test_data = df_prepared.randomSplit([0.8, 0.2], seed=42)

    # Random Forest
    start_time = time.time()
    rf_model = rf.fit(train_data)
    rf_predictions = rf_model.transform(test_data)
    rf_time = time.time() - start_time
    rf_rmse = evaluator_rmse.evaluate(rf_predictions)
    rf_r2 = evaluator_r2.evaluate(rf_predictions)

    # Gradient Boosting
    start_time = time.time()
    gbt_model = gbt.fit(train_data)
    gbt_predictions = gbt_model.transform(test_data)
    gbt_time = time.time() - start_time
    gbt_rmse = evaluator_rmse.evaluate(gbt_predictions)
    gbt_r2 = evaluator_r2.evaluate(gbt_predictions)
```

```

# Store results
results.append((size, rf_time, rf_rmse, rf_r2, gbt_time, gbt_rmse, gbt_r2))
print(f>Data Size: {size} - RF: Time={rf_time:.2f}s, RMSE={rf_rmse:.4f},  

↪R²={rf_r2:.4f} | GBT: Time={gbt_time:.2f}s, RMSE={gbt_rmse:.4f}, R²={gbt_r2:.  

↪4f}")

```

Data Size: 20% - RF: Time=227.44s, RMSE=3.6436, R²=0.9763 | GBT: Time=290.58s, RMSE=1.6169, R²=0.9953

Data Size: 50% - RF: Time=264.25s, RMSE=3.5704, R²=0.9771 | GBT: Time=362.53s, RMSE=1.6424, R²=0.9952

Data Size: 75% - RF: Time=290.90s, RMSE=3.5671, R²=0.9771 | GBT: Time=415.53s, RMSE=1.6325, R²=0.9952

[Stage 5328:> (0 + 1) / 1]

Data Size: 100% - RF: Time=321.87s, RMSE=3.6491, R²=0.9761 | GBT: Time=473.94s, RMSE=1.6164, R²=0.9953

```

[64]: import pandas as pd

# Convert results to DataFrame
results_df = pd.DataFrame(results, columns=["Data Size",
↪"RF Training Time", "RF RMSE", "RF R²",
↪"GBT Training Time", "GBT RMSE",
↪"GBT R²"])

# Display results
print("Scale-In Experiment Results:")
print(results_df)

```

Scale-In Experiment Results:

	Data Size	RF Training Time	RF RMSE	RF R²	GBT Training Time \
0	20%	227.437624	3.643566	0.976288	290.576475
1	50%	264.249860	3.570368	0.977129	362.531611
2	75%	290.900296	3.567099	0.977091	415.527479
3	100%	321.868116	3.649114	0.976130	473.944690

	GBT RMSE	GBT R²
0	1.616928	0.995330
1	1.642350	0.995161

```
2  1.632526  0.995202
3  1.616351  0.995317
```

```
[65]: from tabulate import tabulate

# Display the table with borders
print(tabulate(results_df, headers="keys", tablefmt="grid"))
```

```
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+
|      | Data Size | RF Training Time | RF RMSE | RF R2 | GBT Training
Time | GBT RMSE | GBT R2 |
+====+=====+=====+=====+=====+=====+
=====+=====+=====+
| 0 | 20%      |          227.438 | 3.64357 | 0.976288 |
290.576 | 1.61693 | 0.99533 |
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+
| 1 | 50%      |          264.25  | 3.57037 | 0.977129 |
362.532 | 1.64235 | 0.995161 |
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+
| 2 | 75%      |          290.9   | 3.5671  | 0.977091 |
415.527 | 1.63253 | 0.995202 |
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+
| 3 | 100%     |          321.868 | 3.64911 | 0.97613  |
473.945 | 1.61635 | 0.995317 |
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+
```

```
[66]: from IPython.display import display, HTML

# Render DataFrame as an HTML table with borders
html_table = results_df.to_html(index=False, border=1, justify="center")
display(HTML(html_table))
```

```
<IPython.core.display.HTML object>
```

```
[67]: import matplotlib.pyplot as plt
import numpy as np

# Data for plotting
data_sizes = results_df["Data Size"]
rf_times = results_df["RF Training Time"]
gbt_times = results_df["GBT Training Time"]

# Bar width
```

```

bar_width = 0.35
index = np.arange(len(data_sizes))

# Plot Training Time
plt.figure(figsize=(10, 6))
plt.bar(index, rf_times, bar_width, label="Random Forest", color="skyblue")
plt.bar(index + bar_width, gbt_times, bar_width, label="Gradient Boosting",
        color="lightcoral")

# Add labels and title
plt.xlabel("Data Size")
plt.ylabel("Training Time (seconds)")
plt.title("Training Time Comparison: Random Forest vs Gradient Boosting")
plt.xticks(index + bar_width / 2, data_sizes)
plt.legend()

# Display the plot
plt.tight_layout()
plt.show()

```



```

[68]: # Data for plotting
rf_rmse = results_df["RF RMSE"]
gbt_rmse = results_df["GBT RMSE"]

# Plot RMSE Comparison

```


```

plt.figure(figsize=(10, 6))
plt.plot(data_sizes, rf_rmse, marker="o", label="Random Forest RMSE",
         color="blue", linestyle="--")
plt.plot(data_sizes, gbt_rmse, marker="o", label="Gradient Boosting RMSE",
         color="red", linestyle="--")

# Add labels and title
plt.xlabel("Data Size")
plt.ylabel("RMSE")
plt.title("RMSE Comparison: Random Forest vs Gradient Boosting")
plt.legend()

# Display the plot
plt.tight_layout()
plt.show()

```



```

[69]: # Data for plotting
rf_r2 = results_df["RF R2"]
gbt_r2 = results_df["GBT R2"]

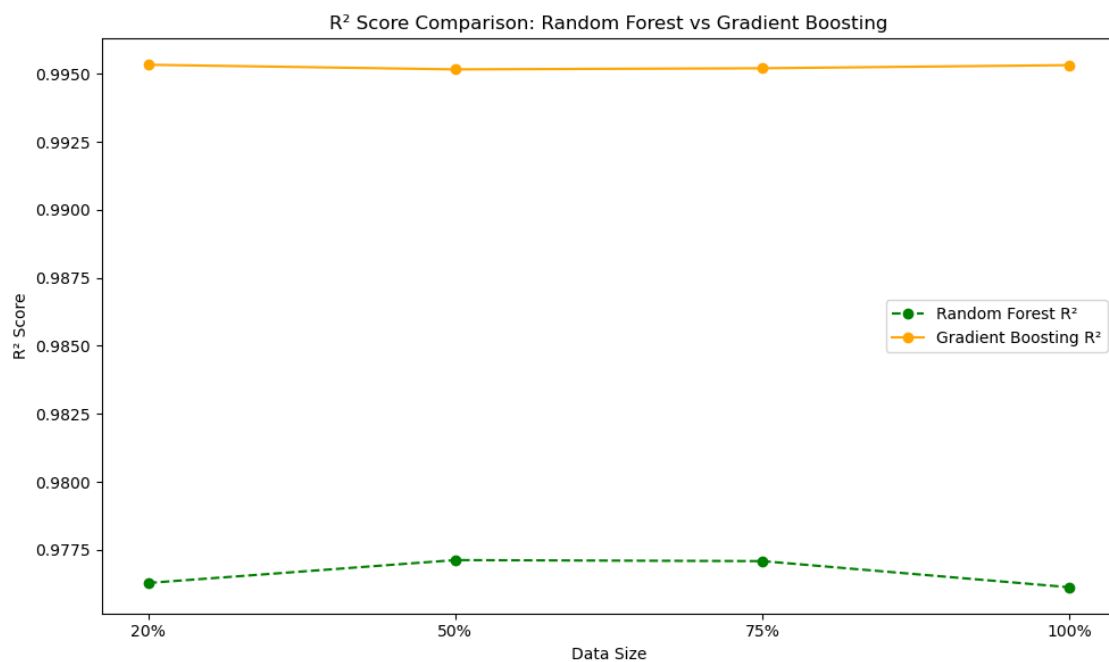
# Plot R2 Comparison
plt.figure(figsize=(10, 6))
plt.plot(data_sizes, rf_r2, marker="o", label="Random Forest R2",
         color="green", linestyle="--")

```

```
plt.plot(data_sizes, gbt_r2, marker="o", label="Gradient Boosting R2",
        color="orange", linestyle="--")

# Add labels and title
plt.xlabel("Data Size")
plt.ylabel("R2 Score")
plt.title("R2 Score Comparison: Random Forest vs Gradient Boosting")
plt.legend()

# Display the plot
plt.tight_layout()
plt.show()
```



```
[70]: # Data for plotting
data_sizes = results_df["Data Size"]
rf_rmse = results_df["RF RMSE"]
gbt_rmse = results_df["GBT RMSE"]
rf_r2 = results_df["RF R2"]
gbt_r2 = results_df["GBT R2"]

# Bar width and positions
bar_width = 0.3
index = np.arange(len(data_sizes))

# Plot RMSE
```



```

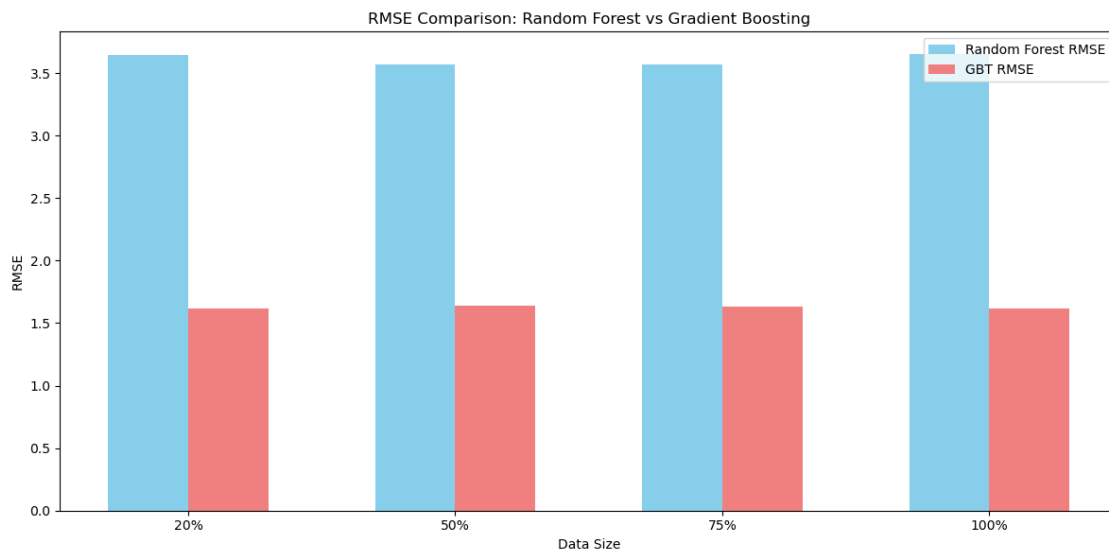
plt.figure(figsize=(12, 6))
plt.bar(index - bar_width/2, rf_rmse, bar_width, label="Random Forest RMSE",
        color="skyblue")
plt.bar(index + bar_width/2, gbt_rmse, bar_width, label="GBT RMSE",
        color="lightcoral")

plt.xlabel("Data Size")
plt.ylabel("RMSE")
plt.title("RMSE Comparison: Random Forest vs Gradient Boosting")
plt.xticks(index, data_sizes)
plt.legend()
plt.tight_layout()
plt.show()

# Plot R2
plt.figure(figsize=(12, 6))
plt.bar(index - bar_width/2, rf_r2, bar_width, label="Random Forest R2",
        color="lightgreen")
plt.bar(index + bar_width/2, gbt_r2, bar_width, label="GBT R2", color="orange")

plt.xlabel("Data Size")
plt.ylabel("R2 Score")
plt.title("R2 Score Comparison: Random Forest vs Gradient Boosting")
plt.xticks(index, data_sizes)
plt.legend()
plt.tight_layout()
plt.show()

```





[]: